

Assignment# 6

Problem 1

In the lab, a classification tree was applied to the Carseats data set after converting Sales into a binary response variable. This question will seek to predict Sales using regression trees and related approaches, treating the response as a quantitative variable (that is, without the conversion).

(a) Split the data set into a training set and a test set.

```
>
> library(ISLR)
> set.seed(1)
> train_data <- sample(1:nrow(Carseats), nrow(Carseats)/2)
> Carseats.train <- Carseats[train_data, ]
> Carseats.test <- Carseats[-train_data, ]
> |
```

(b) Fit a regression tree to the training set. Plot the tree, and interpret the results. Then compute the test MSE.

```
> library(tree)
> tree.carseats <- tree(Sales ~ ., data = Carseats.train)
> summary(tree.carseats)
```

Regression tree:
tree(formula = Sales ~ ., data = Carseats.train)
Variables actually used in tree construction:
[1] "ShelveLoc" "Price" "Age" "Advertising" "Income" "CompPrice"
Number of terminal nodes: 18
Residual mean deviance: 2.36 = 429.5 / 182
Distribution of residuals:

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-4.2570	-1.0360	0.1024	0.0000	0.9301	3.9130

```
> |
```


```
> plot(tree.carseats)
> text(tree.carseats, pretty = 0)
> |
```

>

We can see that MSE is **4.14**

(c) Prune the tree obtained in (b). Use cross validation to determine the optimal level of tree complexity. Plot the pruned tree and interpret the results. Compute the test MSE of the pruned tree. Does pruning improve the test error?


>



Pruning the tree:

In this case, the tree of size 8 is selected by cross-validation. We now prune the tree to obtain the 8-node tree.

```
> prune.carseats <- prune.tree(tree.carseats, best = 8)
> plot(prune.carseats)
> text(prune.carseats, pretty = 0)
> |
```



```

> ycap <- predict(prune.carseats, newdata = Carseats.test)
> mean(ycap - Carseats.test$Sales)^2
[1] 5.09085
> |

```

MSE for the prune tree is 5.09. As we can see pruning the tree has increased the MSE from 4.14 to 5.09.

(d) Use the bagging approach to analyze the data. What test MSE do you obtain? Determine which variables are most important.

```

> bag.carseats <- randomForest(Sales ~ ., data = Carseats.train, mtry = 10, ntree = 500, importance = TRUE)
> ycap.bag <- predict(bag.carseats, newdata = Carseats.test)
> mean(ycap.bag - Carseats.test$Sales)^2
[1] 2.57295
> |

```

We see that bagging decreases the test MSE to **2.57**

To determine which variable is most important we use **importance()** method.

```
> importance(bag.carseats)
              %IncMSE  IncNodePurity
CompPrice    16.297100    130.796172
Income        4.828604     78.208046
Advertising  14.688260    124.933965
Population    2.251331     58.882291
Price        57.016882    517.991476
ShelveLoc    46.096614    319.334615
Age          22.019714    194.098835
Education     2.966678     40.162590
Urban        -2.100855      8.873266
US           7.003729     16.330146
> |
```

From the result, we conclude that price and ShelveLoc are the two most important variables. (due to higher %IncMSE values)

(e) Use random forests to analyze the data. What test MSE do you obtain? Determine which variables are most important.

```
> rf.carseats <- randomForest(Sales ~ ., data = Carseats.train, mtry = 3, ntree = 500, importance = TRUE)
> ycap.rf <- predict(rf.carseats, newdata = Carseats.test)
> mean((ycap.rf - Carseats.test$Sales)^2)
[1] 3.326674
> |
```

In this case, test MSE is **3.32**

```
> importance(rf.carseats)
              %IncMSE  IncNodePurity
CompPrice     8.0728869    130.15824
Income        3.9001054    121.88353
Advertising  12.4449886    138.63178
Population   -0.9082328     97.30322
Price       36.0489662    384.19473
ShelveLoc    31.2271291    242.38514
Age          16.0387820    195.99128
Education     2.1311852     65.66093
Urban        -3.0169350     16.22909
US           5.2815786     33.08008
> |
```

From the results, we see that Price and ShelveLoc are the two most important variables (due to high %IncMSE values).

Problem 2

In the lab, we applied random forests to the Boston data using $mtry=6$ and $ntree=100$.


(a) Consider a more comprehensive range of values for $mtry$: 1, 2,...,13. Given each value of $mtry$, find the test error resulting from random forests on the Boston data (using $ntree=100$). Create a plot displaying the test error rate vs. the value of $mtry$. Comment on the results in the plot.

Creating the test and train data sets and $mtry$ and test error vectors

```
> library(MASS)
> attach(Boston)
> set.seed(1)
> train_data <- sample ( 1: nrow(Boston), nrow(Boston) / 2)
> Boston.train <- Boston[train_data,]
> Boston.test <- Boston[-train_data,]
>
> mtry <- c(1:13)
> test.error <- souble(13)
Error in souble(13) : could not find function "souble"
> test.error <- double(13)
> mtry
[1] 1 2 3 4 5 6 7 8 9 10 11 12 13
> test.error
[1] 0 0 0 0 0 0 0 0 0 0 0 0 0
> |

> for(i in 1:13){
+ rf.Boston <- randomForest(medv ~ ., Boston.train, mtry = i, ntree = 100, importance = T)
+ ycap.rf <- predict(rf.Boston, Boston.test)
+ test.err[i] <- mean((ycap.rf - Boston.test$medv)^2)
+ }
> test.err
[1] 22.94959 14.63641 12.82954 11.85036 11.11480 11.56394 11.53832 11.88346 11.99004 12.42750 13.16399 12.73479
[13] 13.73848
> df <- cbind(mtry, test.err)
> plot(df, col = 'red', type = 'b', pch =19, xlab = 'mtry', ylab = 'Test MSE')
> |
```

Graph Plot:



Looking at the graph, we observe that Test MSE is extremely high when $mtry = 1$ and then decreases as the value of $mtry$ increases. However, with increasing $mtry$ Test MSE again starts to increase.

The lowest Test MSE is when $mtry = 5$.

(b) Similarly, consider a range of values for $ntree$ (between 5 to 200). Given each value of $ntree$, find the test error resulting from random forests (using $mtry=6$). Create a plot displaying the test error vs. the value of $ntree$. Comment on the results in the plot.

Removing the 0.00 values in test.err and creating data frame nTree and Test MSE:

The plot displays the Test Mean Squared Error (MSE) as a function of the number of trees (nTree) in a random forest model. The y-axis, 'Test MSE', is scaled from 11 to 15. The x-axis, 'nTree', is scaled from 0 to 200. The data is represented by a red line with circular markers at each integer value of nTree. The error is highest at nTree=0 (approx. 14.5) and decreases rapidly, reaching a local minimum of approximately 10.7 at nTree=25. Beyond nTree=25, the error continues to decrease but with significant fluctuations, generally staying between 11 and 12.5, indicating that the model's performance stabilizes after a certain number of trees.

From the plot, we observe that when nTree is below 50 Test MSE is on the higher side. However, on increasing the value of nTree Test MSE is reduced with minimum Test MSE coming for nTree value of **70**.